

Assignment 3

Install the Transformers, Datasets, and Evaluate libraries to run this notebook.

```
!pip install datasets evaluate transformers[sentencepiece]
!apt install git-lfs
```

Looking in indexes: <https://pypi.org/simple>, <https://us-python.pkg.dev/colab-wheels/public/simple/>

Collecting datasets
 Downloading datasets-2.8.0-py3-none-any.whl (452 kB)
 452.9/452.9 KB 16.9 MB/s eta 0:00:00

Collecting evaluate
 Downloading evaluate-0.4.0-py3-none-any.whl (81 kB)
 81.4/81.4 KB 11.3 MB/s eta 0:00:00

Collecting transformers[sentencepiece]
 Downloading transformers-4.25.1-py3-none-any.whl (5.8 MB)
 5.8/5.8 MB 63.2 MB/s eta 0:00:00

Requirement already satisfied: pyarrow>=6.0.0 in /usr/local/lib/python3.8/dist-packages (from datasets) (9.0.0)
 Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.8/dist-packages (from datasets) (6.0)
 Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.8/dist-packages (from datasets) (1.21.6)
 Requirement already satisfied: requests>=2.19.0 in /usr/local/lib/python3.8/dist-packages (from datasets) (2.25.1)
 Collecting multiprocessing
 Downloading multiprocessing-0.70.14-py38-none-any.whl (132 kB)
 132.0/132.0 KB 15.2 MB/s eta 0:00:00

Requirement already satisfied: dill<0.3.7 in /usr/local/lib/python3.8/dist-packages (from datasets) (0.3.6)
 Requirement already satisfied: packaging in /usr/local/lib/python3.8/dist-packages (from datasets) (21.3)
 Requirement already satisfied: tqdm>=4.62.1 in /usr/local/lib/python3.8/dist-packages (from datasets) (4.64.1)
 Requirement already satisfied: pandas in /usr/local/lib/python3.8/dist-packages (from datasets) (1.3.5)
 Collecting xxhash
 Downloading xxhash-3.2.0-cp38-cp38-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (213 kB)
 213.0/213.0 KB 19.8 MB/s eta 0:00:00

Collecting responses<0.19
 Downloading responses-0.18.0-py3-none-any.whl (38 kB)
 Requirement already satisfied: aiohttp in /usr/local/lib/python3.8/dist-packages (from datasets) (3.8.3)
 Collecting huggingface-hub<1.0.0,>=0.2.0
 Downloading huggingface-hub-0.11.1-py3-none-any.whl (182 kB)
 182.4/182.4 KB 15.5 MB/s eta 0:00:00

Requirement already satisfied: fsspec[http]>=2021.11.1 in /usr/local/lib/python3.8/dist-packages (from datasets) (2022.11.0)
 Requirement already satisfied: regex!=2019.12.17 in /usr/local/lib/python3.8/dist-packages (from transformers[sentencepiece]) (2022.6)
 Requirement already satisfied: filelock in /usr/local/lib/python3.8/dist-packages (from transformers[sentencepiece]) (3.8.2)
 Collecting tokenizers!=0.11.3,<0.14,>=0.11.1
 Downloading tokenizers-0.13.2-cp38-cp38-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (7.6 MB)
 7.6/7.6 MB 66.2 MB/s eta 0:00:00

Requirement already satisfied: protobuf<=3.20.2 in /usr/local/lib/python3.8/dist-packages (from transformers[sentencepiece]) (3.19.6)
 Collecting sentencepiece!=0.1.92,>=0.1.91
 Downloading sentencepiece-0.1.97-cp38-cp38-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (1.3 MB)
 1.3/1.3 MB 73.4 MB/s eta 0:00:00

Requirement already satisfied: attrs>=17.3.0 in /usr/local/lib/python3.8/dist-packages (from aiohttp->datasets) (22.2.0)
 Requirement already satisfied: yarl<2.0,>=1.0 in /usr/local/lib/python3.8/dist-packages (from aiohttp->datasets) (1.8.2)
 Requirement already satisfied: charset-normalizer<3.0,>=2.0 in /usr/local/lib/python3.8/dist-packages (from aiohttp->datasets) (2.1.1)
 Requirement already satisfied: frozenlist>=1.1.1 in /usr/local/lib/python3.8/dist-packages (from aiohttp->datasets) (1.3.3)
 Requirement already satisfied: multidict<7.0,>=4.5 in /usr/local/lib/python3.8/dist-packages (from aiohttp->datasets) (6.0.3)
 Requirement already satisfied: async-timeout<5.0,>=4.0.0a3 in /usr/local/lib/python3.8/dist-packages (from aiohttp->datasets) (4.0.2)
 Requirement already satisfied: aiosignal>=1.1.2 in /usr/local/lib/python3.8/dist-packages (from aiohttp->datasets) (1.3.1)
 Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/python3.8/dist-packages (from huggingface-hub<1.0.0,>=0.2
 Requirement already satisfied: pyparsing!=3.0.5,>=2.0.2 in /usr/local/lib/python3.8/dist-packages (from packaging->datasets) (3.0.9)
 Requirement already satisfied: chardet<5,>=3.0.2 in /usr/local/lib/python3.8/dist-packages (from requests>=2.19.0->datasets) (4.0.0)
 Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.8/dist-packages (from requests>=2.19.0->datasets) (2022.12.7)
 Requirement already satisfied: idna<3,>=2.5 in /usr/local/lib/python3.8/dist-packages (from requests>=2.19.0->datasets) (2.10)
 Requirement already satisfied: urllib3<1.27,>=1.21.1 in /usr/local/lib/python3.8/dist-packages (from requests>=2.19.0->datasets) (1.26.13)
 Collecting urllib3<1.27,>=1.21.1
 Downloading urllib3-1.26.13-py2.py3-none-any.whl (140 kB)
 140.6/140.6 KB 17.4 MB/s eta 0:00:00

Requirement already satisfied: pytz>=2017.3 in /usr/local/lib/python3.8/dist-packages (from pandas->datasets) (2022.7)

```
from datasets import load_dataset, DatasetDict
```

```
ds_train = load_dataset("huggingface-course/codeparrot-ds-train", split="train")
ds_valid = load_dataset("huggingface-course/codeparrot-ds-valid", split="validation")
```

```
raw_datasets = DatasetDict(
    {
        "train": ds_train.shuffle(seed=42).select(range(50000)),
        "valid": ds_valid.shuffle(seed=42).select(range(500))
    }
)
```

raw_datasets

WARNING:datasets.builder:Using custom data configuration huggingface-course--codep
 Downloading and preparing dataset json/huggingface-course--codeparrot-ds-train to
 Downloading data files: 1/1 [02:48<00:00,
 100% 168.87s/it]
 Downloading data: 8.25G/8.25G [02:48<00:00,
 100% 69.7MB/s]
 WARNING:datasets.download.download_manager:Computing checksums of downloaded files
 Computing checksums: 1/1 [00:45<00:00,
 100% 45.70s/it]
 Extracting data files: 100% 1/1 [00:00<00:00, 38.14it/s]

Dataset json downloaded and prepared to /root/.cache/huggingface/datasets/huggingf
 WARNING:datasets.builder:Using custom data configuration huggingface-course--codep
 Downloading and preparing dataset json/huggingface-course--codeparrot-ds-valid to
 Downloading data files: 1/1 [00:04<00:00,
 100% 4.21s/it]
 Downloading data: 46.1M/46.1M [00:01<00:00,
 100% 41.5MB/s]
 Extracting data files: 100% 1/1 [00:00<00:00, 51.73it/s]

Dataset json downloaded and prepared to /root/.cache/huggingface/datasets/huggingf
 DatasetDict({

```
for key in raw_datasets["train"][0]:
    print(f"{key.upper()}: {raw_datasets['train'][0][key][:200]}")
```

```
REPO_NAME: ThomasMiconi/htmresearch
PATH: projects/feedback/feedback_sequences.py
COPIES: 2
SIZE: 26875
CONTENT:
# Numenta Platform for Intelligent Computing (NuPIC)
# Copyright (C) 2016, Numenta, Inc. Unless you have an agreement
# with Numenta, Inc., for a separate license for this software code, the
# follo
LICENSE: agpl-3.0
```

```
from transformers import AutoTokenizer
```

```
context_length = 128
tokenizer = AutoTokenizer.from_pretrained("huggingface-course/code-search-net-tokenizer")
```

```
outputs = tokenizer(
    raw_datasets["train"][:2]["content"],
    truncation=True,
    max_length=context_length,
    return_overflowing_tokens=True,
    return_length=True,
)
```

```
print(f"Input IDs length: {len(outputs['input_ids'])}")
print(f"Input chunk lengths: {(outputs['length'])}")
print(f"Chunk mapping: {outputs['overflow_to_sample_mapping']}")
```

Downloading: 100% 265/265 [00:00<00:00, 8.34kB/s]
 Downloading: 100% 789k/789k [00:00<00:00, 601kB/s]
 Downloading: 448k/448k [00:00<00:00,
 100% 1.16MB/s]
 Downloading: 1.34M/1.34M [00:00<00:00,
 100% 3.50MB/s]
 Downloading: 100% 90.0/90.0 [00:00<00:00, 2.06kB/s]

Input IDs length: 86

```
def tokenize(element):
    outputs = tokenizer(
        element["content"],
        truncation=True,
        max_length=context_length,
        return_overflowing_tokens=True,
        return_length=True,
    )
    input_batch = []
    for length, input_ids in zip(outputs["length"], outputs["input_ids"]):
        if length == context_length:
            input_batch.append(input_ids)
    return {"input_ids": input_batch}
```

```
tokenized_datasets = raw_datasets.map(
    tokenize, batched=True, remove_columns=raw_datasets["train"].column_names
)
tokenized_datasets
```

100% 50/50 [07:29<00:00, 8.66s/ba]

100% 1/1 [00:04<00:00, 4.43s/ba]

```
DatasetDict({
  train: Dataset({
    features: ['input_ids'],
    num_rows: 1375550
  })
  valid: Dataset({
    features: ['input_ids'],
    num_rows: 13617
  })
})
```

```
from transformers import AutoTokenizer, GPT2LMHeadModel, AutoConfig
```

```
config = AutoConfig.from_pretrained(
    "gpt2",
    vocab_size=len(tokenizer),
    n_ctx=context_length,
    bos_token_id=tokenizer.bos_token_id,
    eos_token_id=tokenizer.eos_token_id,
)
```

Downloading: 100% 665/665 [00:00<00:00, 33.1kB/s]

```
model = GPT2LMHeadModel(config)
model_size = sum(t.numel() for t in model.parameters())
print(f"GPT-2 size: {model_size/1000**2:.1f}M parameters")
```

GPT-2 size: 124.4M parameters

```
from transformers import DataCollatorForLanguageModeling
```

```
tokenizer.pad_token = tokenizer.eos_token
data_collator = DataCollatorForLanguageModeling(tokenizer, mlm=False)
```

```
out = data_collator([tokenized_datasets["train"][i] for i in range(5)])
for key in out:
    print(f"{key} shape: {out[key].shape}")
```

You're using a GPT2TokenizerFast tokenizer. Please note that with a fast tokenizer, using the `__call__` method is faster than using a `encode_plus` call.

```
input_ids shape: torch.Size([5, 128])
attention_mask shape: torch.Size([5, 128])
labels shape: torch.Size([5, 128])
```

▼ Training

Possible Optimizers to try Optimizers = adamw_hf, adamw_torch, adamw_apex_fused, adamw_anyprecision or adafactor.

modify max_steps to stop after a number of iterations

modify batch size to fit into memory modify save every n steps to modify how often save occurs

modify output_dir to a google drive path to save and load the model correctly

```
from transformers import Trainer, TrainingArguments
```

```
args = TrainingArguments(
    output_dir="codeparrot-ds",
    optim= 'adamw_hf',
    per_device_train_batch_size=16,
    per_device_eval_batch_size=16,
    evaluation_strategy="steps",
    eval_steps=5_000,
    logging_steps=1,
    gradient_accumulation_steps=8,
    num_train_epochs=2,
    weight_decay=0.1,
    warmup_steps=100,
    lr_scheduler_type="cosine",
    learning_rate=5e-4,
    save_steps=100,
    fp16=True,
    max_steps=300,
)
trainer = Trainer(
    model=model,
    tokenizer=tokenizer,
    args=args,
    data_collator=data_collator,
    train_dataset=tokenized_datasets["train"],
    eval_dataset=tokenized_datasets["valid"],
)
```

max_steps is given, it will override any value given in num_train_epochs
Using cuda_amp half precision backend

```
result = trainer.train()
```

```
⚠ /usr/local/lib/python3.8/dist-packages/transformers/optimization.py:306: FutureWarning:
  warnings.warn(
```

```
***** Running training *****
Num examples = 1375550
Num Epochs = 1
Instantaneous batch size per device = 16
Total train batch size (w. parallel, distributed & accumulation) = 128
Gradient Accumulation steps = 8
Total optimization steps = 300
Number of trainable parameters = 124439808
```

[300/300 09:29, Epoch 0/1]

Step Training Loss Validation Loss

```
Saving model checkpoint to codeparrot-ds/checkpoint-100
Configuration saved in codeparrot-ds/checkpoint-100/config.json
Model weights saved in codeparrot-ds/checkpoint-100/pytorch_model.bin
tokenizer config file saved in codeparrot-ds/checkpoint-100/tokenizer_config.json
Special tokens file saved in codeparrot-ds/checkpoint-100/special_tokens_map.json
Saving model checkpoint to codeparrot-ds/checkpoint-200
Configuration saved in codeparrot-ds/checkpoint-200/config.json
Model weights saved in codeparrot-ds/checkpoint-200/pytorch_model.bin
tokenizer config file saved in codeparrot-ds/checkpoint-200/tokenizer_config.json
Special tokens file saved in codeparrot-ds/checkpoint-200/special_tokens_map.json
Saving model checkpoint to codeparrot-ds/checkpoint-300
Configuration saved in codeparrot-ds/checkpoint-300/config.json
Model weights saved in codeparrot-ds/checkpoint-300/pytorch_model.bin
tokenizer config file saved in codeparrot-ds/checkpoint-300/tokenizer_config.json
Special tokens file saved in codeparrot-ds/checkpoint-300/special_tokens_map.json
```

Training completed. Do not forget to share your model on huggingface.co/models =)

```
eval_results = trainer.evaluate()
```

```

***** Running Evaluation *****
Num examples = 13617
Batch size = 16

!pip install ml-things
!pip install matplotlib==3.1.3

Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Collecting ml-things
  Downloading ml_things-0.0.1.tar.gz (8.1 MB)
    8.1/8.1 MB 13.8 MB/s eta 0:00:00
  Preparing metadata (setup.py) ... done
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.8/dist-packages (0.24.2)
Requirement already satisfied: numpy in /usr/local/lib/python3.8/dist-packages (1.24.2)
Requirement already satisfied: requests in /usr/local/lib/python3.8/dist-packages (2.31.0)
Requirement already satisfied: tqdm>=4.27 in /usr/local/lib/python3.8/dist-packages (4.64.1)
Collecting ftfy>=5.8
  Downloading ftfy-6.1.1-py3-none-any.whl (53 kB)
    53.1/53.1 KB 4.5 MB/s eta 0:00:00
Collecting matplotlib>=3.4.0
  Downloading matplotlib-3.6.2-cp38-cp38-manylinux_2_12_x86_64.manylinux2010_x86_64.whl (11.5 MB)
    9.4/9.4 MB 8.4 MB/s eta 0:00:00
Requirement already satisfied: wcwidth>=0.2.5 in /usr/local/lib/python3.8/dist-packages (0.2.5)
Requirement already satisfied: cyclor>=0.10 in /usr/local/lib/python3.8/dist-packages (0.10.0)
Collecting contourpy>=1.0.1
  Downloading contourpy-1.0.6-cp38-cp38-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (296.0 KB)
    296.0/296.0 KB 22.6 MB/s eta 0:00:00
Collecting fonttools>=4.22.0
  Downloading fonttools-4.38.0-py3-none-any.whl (965 kB)
    965.4/965.4 KB 50.4 MB/s eta 0:00:00
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.8/dist-packages (1.4.4)
Requirement already satisfied: pyparsing>=2.2.1 in /usr/local/lib/python3.8/dist-packages (3.1.0)
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.8/dist-packages (9.4.0)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.8/dist-packages (23.1)
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.8/dist-packages (2.8.2)
Requirement already satisfied: idna<3, >=2.5 in /usr/local/lib/python3.8/dist-packages (3.4)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.8/dist-packages (2022.12.7)
Requirement already satisfied: chardet<5, >=3.0.2 in /usr/local/lib/python3.8/dist-packages (5.1.0)
Requirement already satisfied: urllib3<1.27, >=1.21.1 in /usr/local/lib/python3.8/dist-packages (1.26.15)
Requirement already satisfied: scipy>=1.1.0 in /usr/local/lib/python3.8/dist-packages (1.10.1)
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.8/dist-packages (3.1.0)
Requirement already satisfied: joblib>=0.11 in /usr/local/lib/python3.8/dist-packages (1.3.2)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.8/dist-packages (1.16.0)
Building wheels for collected packages: ml-things
  Building wheel for ml-things (setup.py) ... done
  Created wheel for ml-things: filename=ml_things-0.0.1-py3-none-any.whl size=24111 sha256=b0137206f860cf08870a4fda0b121a
  Stored in directory: /root/.cache/pip/wheels/b0/13/72/06f860cf08870a4fda0b121a
Successfully built ml-things
Installing collected packages: ftfy, fonttools, contourpy, matplotlib, ml-things
  Attempting uninstall: matplotlib
    Found existing installation: matplotlib 3.2.2
    Uninstalling matplotlib-3.2.2:
      Successfully uninstalled matplotlib-3.2.2
Successfully installed contourpy-1.0.6 fonttools-4.38.0 ftfy-6.1.1 matplotlib-3.6.2 ml-things-0.0.1
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Collecting matplotlib==3.1.3
  Downloading matplotlib-3.1.3-cp38-cp38-manylinux1_x86_64.whl (13.1 MB)
    13.1/13.1 MB 69.7 MB/s eta 0:00:00
Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in /usr/local/lib/python3.8/dist-packages (3.1.0)
Requirement already satisfied: python-dateutil>=2.1 in /usr/local/lib/python3.8/dist-packages (2.8.2)
Requirement already satisfied: cyclor>=0.10 in /usr/local/lib/python3.8/dist-packages (0.10.0)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.8/dist-packages (1.4.4)
Requirement already satisfied: numpy>=1.11 in /usr/local/lib/python3.8/dist-packages (1.24.2)

```

```

from matplotlib import pyplot as plot
import math
from ml_things import plot_dict, fix_text
# Keep track of train and evaluate loss.
loss_history = {'train_loss':[], 'eval_loss':[]}

# Keep track of train and evaluate perplexity.
# This is a metric useful to track for language models.
perplexity_history = {'train_perplexity':[], 'eval_perplexity':[]}

```

```

# Loop through each log history.
for log_history in trainer.state.log_history:

    if 'loss' in log_history.keys():
        # Deal with training loss.
        loss_history['train_loss'].append(log_history['loss'])
        perplexity_history['train_perplexity'].append(math.exp(log_history['loss']))

    elif 'eval_loss' in log_history.keys():
        # Deal with eval loss.
        loss_history['eval_loss'].append(log_history['eval_loss'])
        perplexity_history['eval_perplexity'].append(math.exp(log_history['eval_loss']))


# Plot Losses.
plot_dict(loss_history, start_step=args.logging_steps,
          step_size=args.logging_steps, use_title='Loss',
          use_xlabel='Train Steps', use_ylabel='Values', magnify=2)

print()

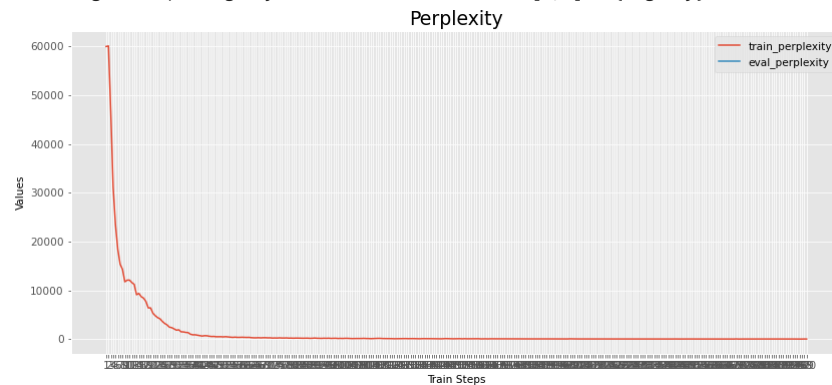
# Plot Perplexities.
plot_dict(perplexity_history, start_step=args.logging_steps,
          step_size=args.logging_steps, use_title='Perplexity',
          use_xlabel='Train Steps', use_ylabel='Values', magnify=2)

```

/usr/local/lib/python3.8/dist-packages/ml_things/plot_functions.py:409: DeprecationWarning: warn(f'`magnify` needs to have value in [0,1]! `{magnify}` will be conv



/usr/local/lib/python3.8/dist-packages/ml_things/plot_functions.py:409: DeprecationWarning: warn(f'`magnify` needs to have value in [0,1]! `{magnify}` will be conv



▼ Report Perplexity and eval_results number with each experiment

```
import numpy as np
print(f"Perplexity: {np.exp(eval_results['eval_loss']):.2f}")
```

Perplexity: 67.44

result

```
TrainOutput(global_step=300, training_loss=5.345289017359415, metrics={'train_runtime': 575.1932, 'train_samples_per_second': 66.76,
'train_steps_per_second': 0.522, 'total_flos': 250840350720000.0, 'train_loss': 5.345289017359415, 'epoch': 0.03})
```

trainer.state.log_history

```
[{'loss': 11.0012, 'learning_rate': 5e-06, 'epoch': 0.0, 'step': 1},
{'loss': 11.0032, 'learning_rate': 1e-05, 'epoch': 0.0, 'step': 2},
{'loss': 10.7478, 'learning_rate': 1.5e-05, 'epoch': 0.0, 'step': 3},
{'loss': 10.3477, 'learning_rate': 2e-05, 'epoch': 0.0, 'step': 4},
{'loss': 10.0628, 'learning_rate': 2.5e-05, 'epoch': 0.0, 'step': 5},
{'loss': 9.8198, 'learning_rate': 3e-05, 'epoch': 0.0, 'step': 6},
{'loss': 9.6403,
 'learning_rate': 3.5000000000000004e-05,
 'epoch': 0.0,
 'step': 7},
{'loss': 9.5701, 'learning_rate': 4e-05, 'epoch': 0.0, 'step': 8},
{'loss': 9.3744,
 'learning_rate': 4.4999999999999996e-05,
 'epoch': 0.0,
 'step': 9},
{'loss': 9.4028, 'learning_rate': 5e-05, 'epoch': 0.0, 'step': 10},
{'loss': 9.4031, 'learning_rate': 5.5e-05, 'epoch': 0.0, 'step': 11},
{'loss': 9.3609, 'learning_rate': 6e-05, 'epoch': 0.0, 'step': 12},
{'loss': 9.328,
 'learning_rate': 6.500000000000001e-05,
 'epoch': 0.0,
 'step': 13},
{'loss': 9.1197,
 'learning_rate': 7.000000000000001e-05,
 'epoch': 0.0,
 'step': 14},
{'loss': 9.1488, 'learning_rate': 7.5e-05, 'epoch': 0.0, 'step': 15},
{'loss': 9.0739, 'learning_rate': 8e-05, 'epoch': 0.0, 'step': 16},
{'loss': 9.0402, 'learning_rate': 8.5e-05, 'epoch': 0.0, 'step': 17},
{'loss': 8.9548,
 'learning_rate': 8.999999999999999e-05,
 'epoch': 0.0,
 'step': 18},
{'loss': 8.7704, 'learning_rate': 9.5e-05, 'epoch': 0.0, 'step': 19},
{'loss': 8.7684, 'learning_rate': 0.0001, 'epoch': 0.0, 'step': 20},
{'loss': 8.5814, 'learning_rate': 0.000105, 'epoch': 0.0, 'step': 21},
{'loss': 8.4831, 'learning_rate': 0.00011, 'epoch': 0.0, 'step': 22},
{'loss': 8.3992, 'learning_rate': 0.000115, 'epoch': 0.0, 'step': 23},
{'loss': 8.3439, 'learning_rate': 0.00012, 'epoch': 0.0, 'step': 24},
{'loss': 8.2087, 'learning_rate': 0.000125, 'epoch': 0.0, 'step': 25},
{'loss': 8.0878,
 'learning_rate': 0.00013000000000000002,
 'epoch': 0.0,
 'step': 26},
{'loss': 7.996, 'learning_rate': 0.000135, 'epoch': 0.0, 'step': 27},
{'loss': 7.8232,
 'learning_rate': 0.00014000000000000001,
 'epoch': 0.0,
 'step': 28},
{'loss': 7.7798, 'learning_rate': 0.000145, 'epoch': 0.0, 'step': 29},
{'loss': 7.6732, 'learning_rate': 0.00015, 'epoch': 0.0, 'step': 30},
{'loss': 7.5354, 'learning_rate': 0.000155, 'epoch': 0.0, 'step': 31},
{'loss': 7.5777, 'learning_rate': 0.00016, 'epoch': 0.0, 'step': 32},
{'loss': 7.3418, 'learning_rate': 0.000165, 'epoch': 0.0, 'step': 33},
{'loss': 7.3266, 'learning_rate': 0.00017, 'epoch': 0.0, 'step': 34},
{'loss': 7.2584, 'learning_rate': 0.000175, 'epoch': 0.0, 'step': 35},
{'loss': 7.2291,
 'learning_rate': 0.00017999999999999998,
```

Example to load from checkpoint Note: move to Drive and get Drive path first

▼ Test Code Prompts

Model and Tokenizer must be present

```

import torch
from transformers import pipeline

device = torch.device("cuda:0") if torch.cuda.is_available() else torch.device("cpu")
print(device)
pipe = pipeline(
    "text-generation",
    model=model,
    tokenizer=tokenizer,
    device=device
)

cuda:0

txt = ""
# create some data
x = np.random.randn(100)
y = np.random.randn(100)

# create scatter plot with x, y
"""
print(pipe(txt, num_return_sequences=1)[0]["generated_text"])

Setting `pad_token_id` to `eos_token_id`:0 for open-end generation.
/usr/local/lib/python3.8/dist-packages/transformers/generation/utils.py:1387: UserWarning: Neither `max_length` nor `max_new_tokens` has
    warnings.warn(
# create some data
x = np.random.randn(100)
y = np.random.randn(100)

# create scatter plot with x, y
X, y = np.zeros((32), dtype=np

<
>

txt = ""
# create some data
x = np.random.randn(100)
y = np.random.randn(100)

# create dataframe from x and y
"""
print(pipe(txt, num_return_sequences=1)[0]["generated_text"])

Setting `pad_token_id` to `eos_token_id`:0 for open-end generation.
# create some data
x = np.random.randn(100)
y = np.random.randn(100)

# create dataframe from x and y
X = [[1.5, 2, 2]

Y_

txt = ""
# dataframe with profession, income and name
df = pd.DataFrame({'profession': x, 'income':y, 'name': z})

# calculate the mean income per profession
"""
print(pipe(txt, num_return_sequences=1)[0]["generated_text"])

Setting `pad_token_id` to `eos_token_id`:0 for open-end generation.
# dataframe with profession, income and name
df = pd.DataFrame({'profession': x, 'income':y, 'name': z})

# calculate the mean income per profession
# = self.

txt = ""
# import random forest regressor from scikit-learn
from sklearn.ensemble import RandomForestRegressor

# fit random forest model with 300 estimators on X, y:
"""
print(pipe(txt, num_return_sequences=1)[0]["generated_text"])

```



```
Setting `pad_token_id` to `eos_token_id`:0 for open-end generation.
```

```
# import random forest regressor from scikit-learn  
from sklearn.ensemble import RandomForestRegressor
```

```
# fit random forest model with 300 estimators on X, y:  
import numpy.random.RandomState(0, 20]))  
clf = Grid
```

